

SUPERVISED MACHINE LEARNING: REGRESSION

MODULE1:

INTRODUCTION TO SUPERVISED MACHINE LEARNING AND LINEAR REGRESSION

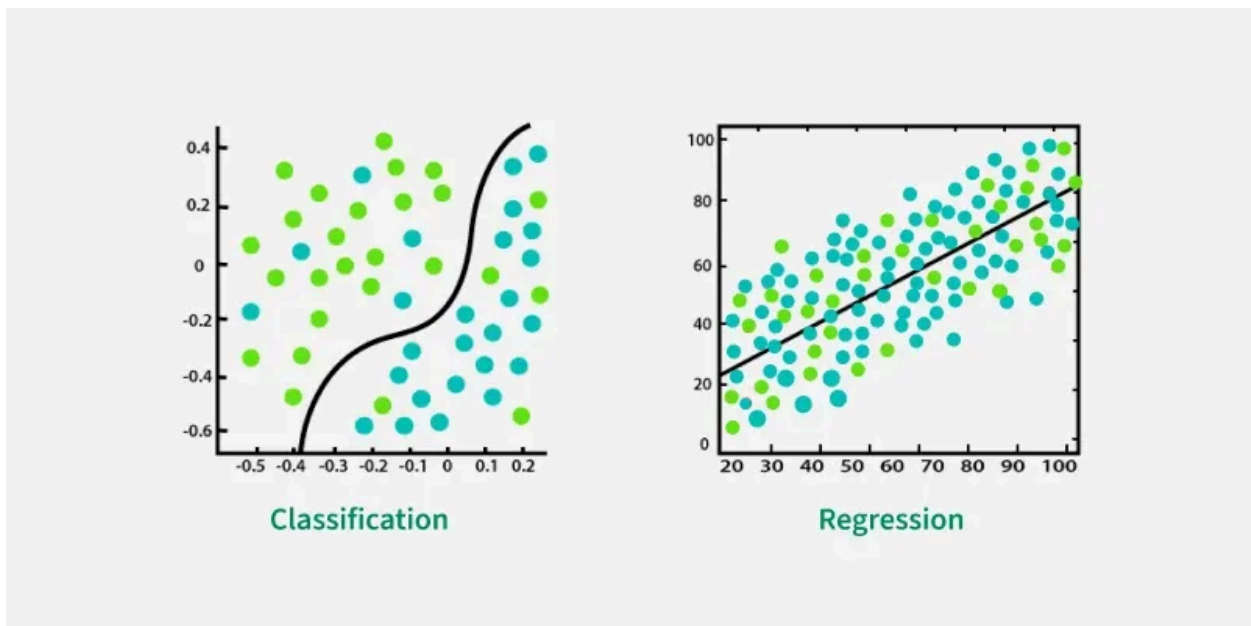
TABLE OF CONTENTS

Introduction to Supervised Machine Learning.....	2
Models and Parameters.....	2
Training a Model.....	3
Trade-off: Interpretation vs Prediction.....	4
Linear Regression.....	4
Implementation in Python (scikit-learn).....	5
Data Transformations.....	6
Key Takeaways.....	8

Introduction to Supervised Machine Learning

- **Machine Learning (ML)** enables computers to learn from data instead of being explicitly programmed.
- Unlike traditional statistical models (where we often know the data-generating process), ML deals with complex or unknown processes.
- ML is a subset of **Artificial Intelligence (AI)**, focusing on **learning ability**.

Two main types of Supervised Learning:



- **Regression:** predicts a **continuous variable** (e.g., house prices, movie revenue).
- **Classification:** predicts a **categorical variable** (e.g., churn, spam detection, face recognition).

Models and Parameters

- **Model:** a simplified representation of reality that captures important relationships.

Concept	Definition	Example
Variable	A symbol or placeholder representing data values.	x_1, x_2, y
Coefficient	A numeric value that represents the weight/strength of a feature's impact on the target.	β_1, β_2
Argument	A concrete value passed into a function or model.	$f(3) \rightarrow 3$ is argument
Parameter	Quantities the model learns from data, define how the model maps inputs to outputs.	coefficients β
Hyperparameter	Settings chosen before training, which control the model's structure or learning process.	<code>learning_rate=0.01, epochs=50</code>

Training a Model

1. Collect labeled data (X, y) .
2. Split into **training** and **testing** sets to prevent overfitting.
3. Train the model by finding parameters that minimize the **loss function**

General formula:

$$y_p = f(\Omega, x)$$

where Ω are learned parameters.

- **Loss function:** $J(y, \hat{y})$ measures the error between predictions and actual values.
- Example (linear regression):

$$J(\beta_0, \beta_1) = \frac{1}{2m} \sum_{i=1}^m \left((\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$

Trade-off: Interpretation vs Prediction

- **Interpretation:** explain model + feature impact (sales drivers, safety, marketing).
- **Prediction:** maximize accuracy, accept black-box models (churn, defaults, purchases).
- **Trade-off:** balance depends on business goals; both can complement each other.
- **Examples:**
 - Housing prices → interpret feature importance vs predict prices.
 - Customer churn → understand reasons vs predict who leaves.

Linear Regression

- Model:

$$y_{\beta}(x) = \beta_0 + \beta_1 x$$

- Sum of Square Error (SSE):

$$\sum_{i=1}^m \left(y_{\beta}(x^{(i)}) - y_{obs}^{(i)} \right)^2$$

- Total Sum of Squares (TSS):

$$\sum_{i=1}^m \left(\overline{y_{obs}} - y_{obs}^{(i)} \right)^2$$

- Coefficient of Determination (R2):

$$1 - \frac{SSE}{TSS}$$

Implementation in Python (scikit-learn)

- Import libraries

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score
```

- Split train and test set

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)
```

- Train model

```
lr = LinearRegression()
lr.fit(X_train, y_train)
```

- Predict

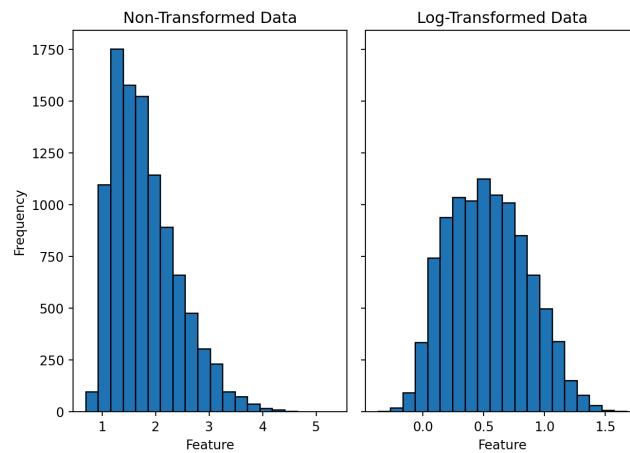
```
y_pred = lr.predict(X_test)
```

- Evaluate

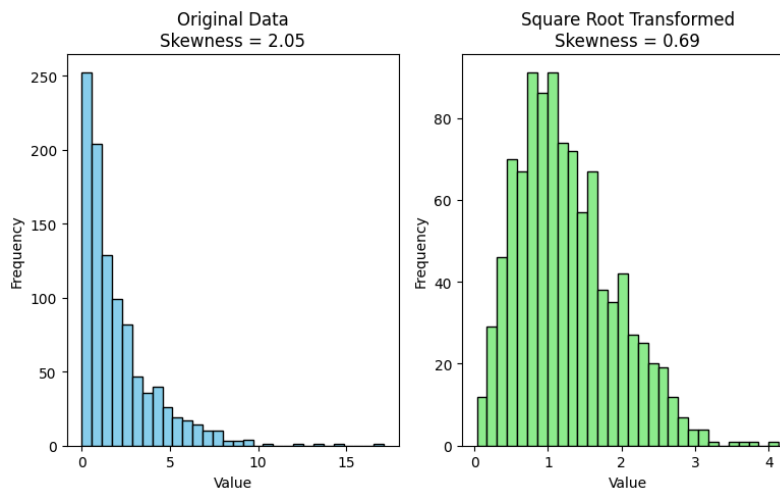
```
print("R2:", r2_score(y_test, y_pred))
```

Data Transformations


- To improve target distribution, common transformations include:
 - **Log transform:** $y' = \log(y+1)$



- **Square root transform:** $y' = \sqrt{y+c}$



- **Box-Cox transform** (chooses optimal λ automatically)



- Example with Box-Cox:

- Import libraries

```
from scipy.stats import boxcox
from scipy.special import inv_boxcox
```

- Transform target variable

```
y_train_bc, lam = boxcox(y_train)
```

- Inverse transform predictions

```
y_pred = inv_boxcox(y_pred_bc, lam)
```

Key Takeaways

- **Machine Learning:** a subset of AI that learns from data.
- **Supervised Learning:**
 - **Regression** → predicts numbers (continuous).
 - **Classification** → predicts categories.
- **Concept:**
 - Variable
 - Coefficient
 - Argument
 - Parameter
 - Hyperparameter
- **Train/Test Split:** needed to avoid overfitting.
- **Interpretation vs Prediction:** trade-off between explainability and accuracy.
- **Linear Regression:** basic supervised model, minimizes MSE.
- **R² Score:** measures how much variance the model explains.
- **Transformations:** log, sqrt, Box-Cox.
- **Python (scikit-learn):** simple tools for training, predicting, evaluating.